

CS578 Statistical Machine Learning Lecture 6

Jean Honorio
Purdue University

(based on slides by Tommi Jaakkola, MIT CSAIL)

Today's topics

- Rating (ordinal regression)
 - reduction to binary problems
 - SVM solution, on-line solution
- Ranking
 - ranking SVM

Rating problems


- A common prediction problem in recommender systems involves rating items (movies, products) on the basis some known features about such objects
- The rating scale is often 1-5 stars assigned to the object
- The key difference between rating problems and multi-way classification problems is that the rating scale is ordinal (e.g., $1 < 2 < 3 < 4 < 5$) while class labels in multi-way classification problems are category symbols

Ordinal regression: setup

- Each item x_i is associated with a feature vector $\phi(x_i)$
 - e.g., product description, movie features, etc.
- We wish to predict an ordinal label $y_i \in \{1, \dots, k\}$ for each item (reflecting views of one user)
- As in the multi-class setting, we translate each rating into a set of binary labels





Binary translation

- There are many ways to translate ratings into binary labels...



Binary translation

- There are many ways to translate ratings into binary labels...







		binary tasks			
rating y	1	-1	0	0	0
	2	1	-1	0	0
	3	1	1	-1	0
	4	1	1	1	-1
	5	1	1	1	1

These are independent partitions that do not enforce the ordinal scale

Binary translation

- We can create more relevant partitions by “sliding” across the ordinal scale



rating





y

binary tasks

1	-1	-1	-1	-1
2	1	-1	-1	-1
3	1	1	-1	-1
4	1	1	1	-1
5	1	1	1	1

Binary translation

- We can create more relevant partitions by “sliding” across the ordinal scale








		binary tasks			
rating y	1	-1	-1	-1	-1
	2	1	-1	-1	-1
	3	1	1	-1	-1
	4	1	1	1	-1
	5	1	1	1	1

The partitions are now dependent ... need to enforce consistency of the binary labels across partitions

Ordinal regression

- We can specify a set of classifiers with shared parameters that always produce consistent binary labels



$$b_1 \leq b_2 \leq b_3 \leq b_4$$

Ordinal regression

- We can specify a set of classifiers with shared parameters that always produce consistent binary labels

$$b_1 \leq b_2 \leq b_3 \leq b_4$$

thresholds are
different but
ordered

$$\begin{array}{c} - \quad + \\ \diagdown \quad \diagup \\ \{1\} \quad \{2,3,4,5\} \end{array} = s_1 = \text{sign}(\underline{\theta \cdot \phi(x_i)} - b_1)$$

common prediction


$$\begin{array}{c} - \quad + \\ \diagdown \quad \diagup \\ \{1,2\} \quad \{3,4,5\} \end{array} = s_2 = \text{sign}(\underline{\theta \cdot \phi(x_i)} - b_2)$$

$$\begin{array}{c} - \quad + \\ \diagdown \quad \diagup \\ \{1,2,3\} \quad \{4,5\} \end{array} = s_3 = \text{sign}(\underline{\theta \cdot \phi(x_i)} - b_3)$$

$$\begin{array}{c} - \quad + \\ \diagdown \quad \diagup \\ \{1,2,3,4\} \quad \{5\} \end{array} = s_4 = \text{sign}(\underline{\theta \cdot \phi(x_i)} - b_4)$$

Ordinal regression, 2nd view


- Each item x_i is associated with a feature vector $\phi(x_i)$
 - e.g., product description, movie features, etc.
- We wish to predict an ordinal label $y_i \in \{1, \dots, k\}$ for each item (reflecting views of one user)
- We assume that there exists an underlying continuous scale from which ratings are obtained via thresholding



$$b_1 \leq b_2 \leq b_3 \leq b_4$$

Ordinal regression, 2nd view

- Each item x_i is associated with a feature vector $\phi(x_i)$
 - e.g., product description, movie features, etc.
- We wish to predict an ordinal label $y_i \in \{1, \dots, k\}$ for each item (reflecting views of one user)
- We assume that there exists an underlying continuous scale from which ratings are obtained via thresholding




Ordinal regression, SVM style

- Given a training set $D = \{(x_i, y_i)\}_{i=1, \dots, n}$
 minimize $\frac{1}{2} \|\underline{\theta}\|^2$ with respect to $\underline{\theta}, b_1, \dots, b_{k-1}$
 such that $b_1 \leq b_2 \leq \dots \leq b_{k-1}$ and
 $s_{il}(\underline{\theta} \cdot \underline{\phi}(x_i) - b_l) \geq 1, \quad l = 1, \dots, k-1, \quad i = 1, \dots, n$


binary classification
constraint

k-1 binary labels
obtained from each
observed rating



Ordinal regression, SVM style

- Given a training set $D = \{(x_i, y_i)\}_{i=1, \dots, n}$
- For instance, assume $k = 5$



- For sample i :

	$l = 1$	$l = 2$	$l = 3$	$l = 4$
$y_i = 3$	$s_{il} = +1$	$s_{il} = +1$	$s_{il} = -1$	$s_{il} = -1$
$y_i = 1$	$s_{il} = -1$	$s_{il} = -1$	$s_{il} = -1$	$s_{il} = -1$

Ordinal regression, PRank

- We can also define a mistake driven perceptron algorithm for solving ordinal regression problems
- The updates are modified slightly due to shared parameters

cycle through the training set $i = 1, \dots, n$

for each example i

$$E_i = \{l : s_{il}(\underline{\theta} \cdot \underline{\phi}(x_i) - b_l) \leq 0\} \quad \text{identify all binary mistakes}$$

$$\underline{\theta} \leftarrow \underline{\theta} + \left(\sum_{l \in E_i} s_{il} \right) \underline{\phi}(x_i) \quad \text{perform a collective update based on the mistakes}$$

$$b_l \leftarrow b_l - s_{il}, \quad l \in E_i \quad \text{update thresholds of each classifier}$$

Note: having a threshold is equivalent to having an extra feature, in which all samples have -1. Thus, the update rule for b_l is not surprising.

Ordinal regression, PRank

- We can also define a mistake driven perceptron algorithm for solving ordinal regression problems
- The updates are modified slightly due to shared parameters

cycle through the training set $i = 1, \dots, n$

for each example i

$$E_i = \{l : s_{il}(\underline{\theta} \cdot \underline{\phi}(x_i) - b_l) \leq 0\} \quad \text{identify all binary mistakes}$$

$$\underline{\theta} \leftarrow \underline{\theta} + \left(\sum_{l \in E_i} s_{il} \right) \underline{\phi}(x_i) \quad \text{perform a collective update based on the mistakes}$$

$$b_l \leftarrow b_l - s_{il}, \quad l \in E_i \quad \text{update thresholds of each classifier}$$

- **Lemma:** if the thresholds are set to zero initially, they will maintain the correct ordering in the course of the algorithm

(See Lemma 1 in [1] if interested in the proof.)

PRank, mistake bound

- **Theorem:** Assume that there exists $\underline{\theta}^*, b_1^*, \dots, b_{k-1}^*$

$$\|\underline{\theta}^*\|^2 + \sum_{l=1}^{k-1} b_l^{*2} = 1$$

such that

$$s_{il}(\underline{\theta}^* \cdot \phi(x_i) - b_l^*) \geq \gamma, \quad l = 1, \dots, k-1, \quad i = 1, \dots, n$$

then the algorithm makes at most

$$(k-1) \frac{R^2 + 1}{\gamma^2}$$

binary mistakes on the training set.

(See Theorem 2 in [1] if interested in the proof.)

Today's topics

- Rating (ordinal regression)
 - reduction to binary problems
 - SVM solution, on-line solution
- Ranking
 - ranking SVM

Ranking

- Rating products, movies, etc. using a few values (e.g., 1-5 stars) results in a partial ranking of the items
- Many rating / classification problems are better viewed as ranking problems
 - suggest movies in the order of user interest in them,
 - rank websites to display in response to a query,
 - suggest genes relevant to a particular disease condition, etc.
- By casting the learning problem as a ranking problem we can also incorporate other types of data / feedback
 - e.g., click through data from users

Ranking example

- We would like to rank n websites (find top sites to display) in response to a few query words

x = context (set of query words)

y = website

(x_1, y_2)

(x_1, y_{10})

(x_1, y_3)

...

(x_1, y_n)

(x_2, y_7)

(x_2, y_2)

(x_2, y_1)

...

(x_2, y_4)

$x_1 = \{ \text{ranking applications} \}$ $x_2 = \{ \text{ranking SVM code} \}$

Ranking example

- We would like to rank n websites (find top sites to display) in response to a few query words

x = context (set of query words)

y = website

(x_1, y_2)

(x_1, y_{10})

(x_1, y_3) **x**

...

(x_1, y_n)

(x_2, y_7) **x**

(x_2, y_2)

(x_2, y_1)

...

(x_2, y_4)

$x_1 = \{ \text{ranking applications} \}$ $x_2 = \{ \text{ranking SVM code} \}$

- The available data contain user selections (clicks) of websites out of those displayed to them

From selections to preferences

- We can interpret a user click as a statement that they prefer the selected link over others displayed in the context of the query

(x_1, y_2)
 (x_1, y_{10})
 (x_1, y_3) **x**

...

(x_1, y_n)

$x_1 = \{ \text{ranking applications} \}$



$(x_1, y_3) > \{ (x_1, y_{10}), (x_1, y_2) \}$

(x_2, y_7) **x**
 (x_2, y_2)
 (x_2, y_1)

...

(x_2, y_4)

$x_2 = \{ \text{ranking SVM code} \}$



$(x_2, y_7) > \{ (x_2, y_2), (x_2, y_1) \}$

Ranking function

- Our goal is to estimate a ranking function over pairs $f(x,y)$ such that its values are consistent with the observed preferences.

$$(x_2, y_7) > \{(x_2, y_2), (x_2, y_1)\}$$

$$\Rightarrow f(x_2, y_7) > f(x_2, y_2), \quad f(x_2, y_7) > f(x_2, y_1)$$

- We can parameterize this function in terms of feature vectors extracted from each pair (context, website)

$$f(x, y; \underline{\theta}) = \underline{\theta} \cdot \underline{\phi}(x, y)$$

Ranking function

- Our goal is to estimate a ranking function over pairs $f(x,y)$ such that its values are consistent with the observed preferences.

$$(x_2, y_7) > \{(x_2, y_2), (x_2, y_1)\}$$

$$\Rightarrow f(x_2, y_7) > f(x_2, y_2), \quad f(x_2, y_7) > f(x_2, y_1)$$

- We can parameterize this function in terms of feature vectors extracted from each pair (context,website)

$$f(x, y; \underline{\theta}) = \underline{\theta} \cdot \underline{\phi}(x, y)$$


where the features could be, e.g.,

$$\phi_w(x, y) = \left\{ \begin{array}{ll} 1, & \text{if word } w \text{ appears in } x \text{ and } y \\ 0, & \text{otherwise} \end{array} \right\}$$

for all $w \in \mathcal{W}$

Ranking function

- The ranking function gives rise to a total ordering of the pairs via projection to the parameter vector



$$f(x, y; \underline{\theta}) = \underline{\theta} \cdot \underline{\phi}(x, y)$$

SVM rank

- A training set of order relations between pairs

$$D = \{ \{ (x_i, y_j) > (x_k, y_l) \} \}$$

- An SVM style algorithm for finding a consistent ranking function

minimize $\frac{1}{2} \|\underline{\theta}\|^2$ with respect to $\underline{\theta}$ such that

$$\underline{\theta} \cdot \underline{\phi}(x_i, y_j) \geq \underline{\theta} \cdot \underline{\phi}(x_k, y_l) + 1,$$

$$\forall \{ (x_i, y_j) > (x_k, y_l) \} \text{ in } D$$

SVM rank

- A training set of order relations between pairs

$$D = \{ \{ (x_i, y_j) > (x_k, y_l) \} \}$$


- An SVM style algorithm for finding a consistent ranking function

$$\text{minimize } \frac{1}{2} \|\underline{\theta}\|^2 + C \sum_{ij;kl} \xi_{ij;kl} \text{ subject to}$$

$$\underline{\theta} \cdot \underline{\phi}(x_i, y_j) \geq \underline{\theta} \cdot \underline{\phi}(x_k, y_l) + 1 - \xi_{ij;kl}, \quad \xi_{ij;kl} \geq 0 \\ \forall \{ (x_i, y_j) > (x_k, y_l) \} \text{ in } D$$

- It is important to appropriately weight or choose which constraints to include

The effect of ranking constraints



adding a single constraint
can have a large effect on
the ranking solution

